

UIMA Tutorial – uimaFIT & DKPro Core



TECHNISCHE
UNIVERSITÄT
DARMSTADT

3rd UIMA@GSCL Workshop, GSCL 2013, Darmstadt

This work is licensed under a
[Attribution-NonCommercial-ShareAlike 4.0 International](https://creativecommons.org/licenses/by-nc-sa/4.0/).
If you are interested in using this material under different conditions, please contact us.



Agenda



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- What is a pipeline?
- Working with annotations
 - What is a type system?
 - What is the Common Analysis Structure (CAS)?
- Working with components
 - What is a reader?
 - What is an analysis engine?
 - What is a writer? (aka consumer)
- DKPro Core component collection

\$?†?‡→○△%† \$%\$?▷⊥ →○▷↑ ♦▷%¥ ⊥↑%▷↑ -- \$▷.:▽▷▷ ↑♦▷} #%‡ ↑%†※
 ▷※▷‡ \$▷.:▽▷▷ ↑♦▷} #%‡ \$%\$†▷ % ⊥?‡→○†▷ ▯▷¥, \$%\$?▷⊥ ?‡
 \$?†?‡→○△%† ♦▽△⊥▷♦▽†¥⊥ =%] →○▷↑ % ♦▷%¥ ⊥↑%▷↑ ?‡†?.:▷, %##▽▷
 ¥?‡→○ ↑▽ % ↑▷%= ▽.: ⊥#?▷‡↑?⊥↑⊥ ?‡ ?↑%†}. ▷%↑♦▷▷ ↑♦%‡
 #▽‡.:△⊥?‡→○ \$%\$?▷⊥, ♦▷%▷?‡→○ =▷▷▷ ↑♦%‡ ▽‡▷ †%‡→○△%→○▷
 →○?※▷⊥ ‡▷▷\$▷▷‡⊥ % =▷‡↑%† \$▽▽⊥↑, %##▽▷¥?‡→○ ↑▽ ↑♦▷ ‡▷▷ ⊥↑△¥],
 ▯♦?#♦ ↑▷⊥↑▷¥ ⊥▷※▷‡=-▷‡↑♦-▽†¥ ?‡.:%‡↑⊥ "?‡‡=%‡} ▷△▷▷▷%‡
 #▽△‡↑▷?▷⊥, .%▷▷‡↑⊥ %▷▷ ▯%▷} ▽.: →○?※?‡→○ % \$?†?‡→○△%† ▷¥△#
 %↑?▽‡ ↑▽ ↑♦?▷※ ?¥⊥ %‡¥ ↑▷} ↑▽ ⊥▷%※ ▽‡†} ▽‡▷ †%‡→○△%→○▷," ⊥
 %?¥ ⊥↑△¥] %△↑♦▷▷ ↗%#▷△⊥ =▷♦†▷▷ ▽.: ↑♦▷ †%‡→○△%→○▷,
 #▽→○‡?↑?▽‡, %‡¥ ¥▷※▷†▽.=▷‡↑ †%\$ %↑ ↑♦▷ ?‡↑▷▷‡%↑?▽‡%† ⊥#♦▽▽†
 .:▷▷ %¥※%‡#▷¥ ⊥↑△¥?▷⊥ ?‡ ↑▷?▷⊥↑▷, ?↑%†}. "↑♦▷} %▷▷ %.:▷%?¥
 [↑♦▷?▷ #♦?†¥▷▷‡] =?→○♦↑ ⊥△.:.:▷▷ ▯▷‡ ↑♦▷} →○▷↑ ↑▽ ⊥#♦▽▽† %‡¥ ⊥▽
 ▽‡," =▷♦†▷▷ ⊥%?¥. "\$▷#%△⊥▷ ▽.: ▽△▷ ▷▷⊥△†↑⊥, ? ¥▽△\$↑ ↑♦%↑ ※▷▷}
 =△#♦."

What is UIMA?



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Component-based architecture for analysis of unstructured data
- „Analysis“ means deriving a structure from the unstructured data
- How does it work?

Like an assembly line...

Take the raw material

Refine it step by step

Drive off with a nice car

Output Example (UIMA Annotation Editor)



TECHNISCHE
UNIVERSITÄT
DARMSTADT

The screenshot displays the UIMA Annotation Editor interface. The main window shows a text document titled 'wikipedia.txt.xml' with the following content:

Wikipedia (i/ˌwɪkiˈpiːdiə/ or i/ˌwɪkiˈpiːdiə/ wik-i-pee-dee-ə) is a free, collaboratively edited and multilingual Internet encyclopedia supported by the non-profit Wikimedia Foundation. Its 22 million articles (over 3.9 million in English alone) have been written collaboratively by volunteers around the world. Almost all of its articles can be edited by anyone with access to the site,[3] and it has about 100,000 regularly active contributors.[4] As of May 2012, there are editions of Wikipedia in 285 languages. It has become the largest and most popular general reference work on the Internet, [5][6][7][8] ranking sixth globally among all websites on Alexa and having an estimated 365 million readers worldwide.[5][9] It is estimated that Wikipedia receives 2.7 billion monthly pageviews from the United States alone.[10]

Wikipedia was launched in January 2001 by Jimmy Wales and Larry Sanger.[11] Sanger coined the name Wikipedia,[12] which is a portmanteau of wiki (a type of collaborative website, from the Hawaiian word wiki, meaning "quick")[13] and encyclopedia. Wikipedia's departure from the expert-driven style of encyclopedia building and the presence of a large body of unacademic content have received extensive attention in print media. In its 2006 Person of the Year article, Time magazine recognized the rapid growth of online collaboration and interaction by millions of people around the world. It cited Wikipedia as an example, in addition to YouTube, MySpace, and Facebook.[14] Wikipedia has also

The interface includes several panels:

- Outline:** A tree view showing the document structure with nodes like Text, ADJ, ADV, NN, etc.
- FeatureStructure View:** A panel for viewing and editing feature structures.
- Annotation Styles:** A table for defining annotation styles.
- Edit View:** Two panels for editing annotations, showing feature and value columns.

Type	Style
<input type="checkbox"/> QUANTMOD	BACKGROUND
<input type="checkbox"/> RCMOD	BACKGROUND
<input type="checkbox"/> REF	BACKGROUND
<input type="checkbox"/> REL	BACKGROUND
<input type="checkbox"/> TMOD	BACKGROUND
<input type="checkbox"/> XCOMP	BACKGROUND
<input type="checkbox"/> XSUBJ	BACKGROUND
<input checked="" type="checkbox"/> ADJ	BACKGROUND
<input checked="" type="checkbox"/> ADV	BACKGROUND
<input type="checkbox"/> ART	BACKGROUND
<input type="checkbox"/> CARD	BACKGROUND
<input type="checkbox"/> CONJ	BACKGROUND
<input type="checkbox"/> N	BACKGROUND
<input checked="" type="checkbox"/> NN	BACKGROUND
<input checked="" type="checkbox"/> NP	BACKGROUND
<input type="checkbox"/> O	BACKGROUND
<input type="checkbox"/> AT	BACKGROUND
<input type="checkbox"/> DM	BACKGROUND
<input type="checkbox"/> EMO	BACKGROUND
<input type="checkbox"/> HASH	BACKGROUND
<input type="checkbox"/> INT	BACKGROUND
<input type="checkbox"/> NNV	BACKGROUND

Feature	Value
sofa	[Sofa]
begin	0
end	185
Governor	[Token]
Dependent	[Token]
Dependency...	punct

Apache UIMA™ - Some history

Unstructured Information Management Architecture



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- 2003 – David Ferrucci and Adam Lally paper
*Accelerating corporate research in the development,
application and deployment of human language
technologies*
- 2004 – IBM alphaWorks project
 - still used e.g. in IBM LanguageWare
- 2006 – Apache Incubator project
- 2009 – OASIS Standard
- 2010 – Full Apache project
- 2010 – Used in IBM's *Watson*
Jeopardy Challenge

<http://uima.apache.org>

Important features of UIMA



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Common representation format for annotations
- Customizable annotation type system
- Common API for analysis components
- Focus on the ability to scale
 - Past: UIMA-CPE – Collection Processing Engine
 - Present: UIMA-AS – Asynchronous Scale-out
 - Future: UIMA-DUCC – Distributed UIMA Cluster Computing
- Active community

- Create and configure pipelines easily in Java
- Test UIMA components
- Started out as a collaborative effort between
 - Center for Computational Pharmacology, University of Colorado, Denver
 - Center for Computational Language and Education Research, University of Colorado, Boulder,
 - Ubiquitous Knowledge Processing (UKP) Lab, Technische Universität Darmstadt
- Since version 2.0.0 part of the Apache UIMA project

- Philip V. Ogren, Steven J. Bethard (2009) **Building Test Suites for UIMA Components**. Proceedings of the Workshop on Software Engineering, Testing, and Quality Assurance for Natural Language Processing (SETQA-NLP 2009). June 2009.
- Christophe Roeder, Philip V. Ogren, William A. Baumgartner Jr., Lawrence Hunter (2009). **Simplifying UIMA Component Development and Testing with Java Annotations and Dependency Injection**, in Chiarcos, C., Eckhart de Castilho, Stede, M. (eds), Von der Form zur Bedeutung: Text automatisch verarbeiten / From Form to Meaning: Processing Texts Automatically. Tübingen: Narr, 2009

<http://uima.apache.org/uimafit.html>

Important features of uimaFIT



TECHNISCHE
UNIVERSITÄT
DARMSTADT

uimaFIT is key to make UIMA usable within Java code

- **Factories – dynamic assembly of analysis pipelines**
 - Automatic type-system detection
 - Most metadata maintained in Java
 - Refactorable code
- **Injection – convenient implementation of analysis components**
 - Default parameter values
 - Parameter types not supported by UIMA (e.g. File, URL, ...)
- **Testing – easy running of analysis pipelines**
 - Unit tests easy to set up
 - ... or research experiments

... and more ...

Pipelines

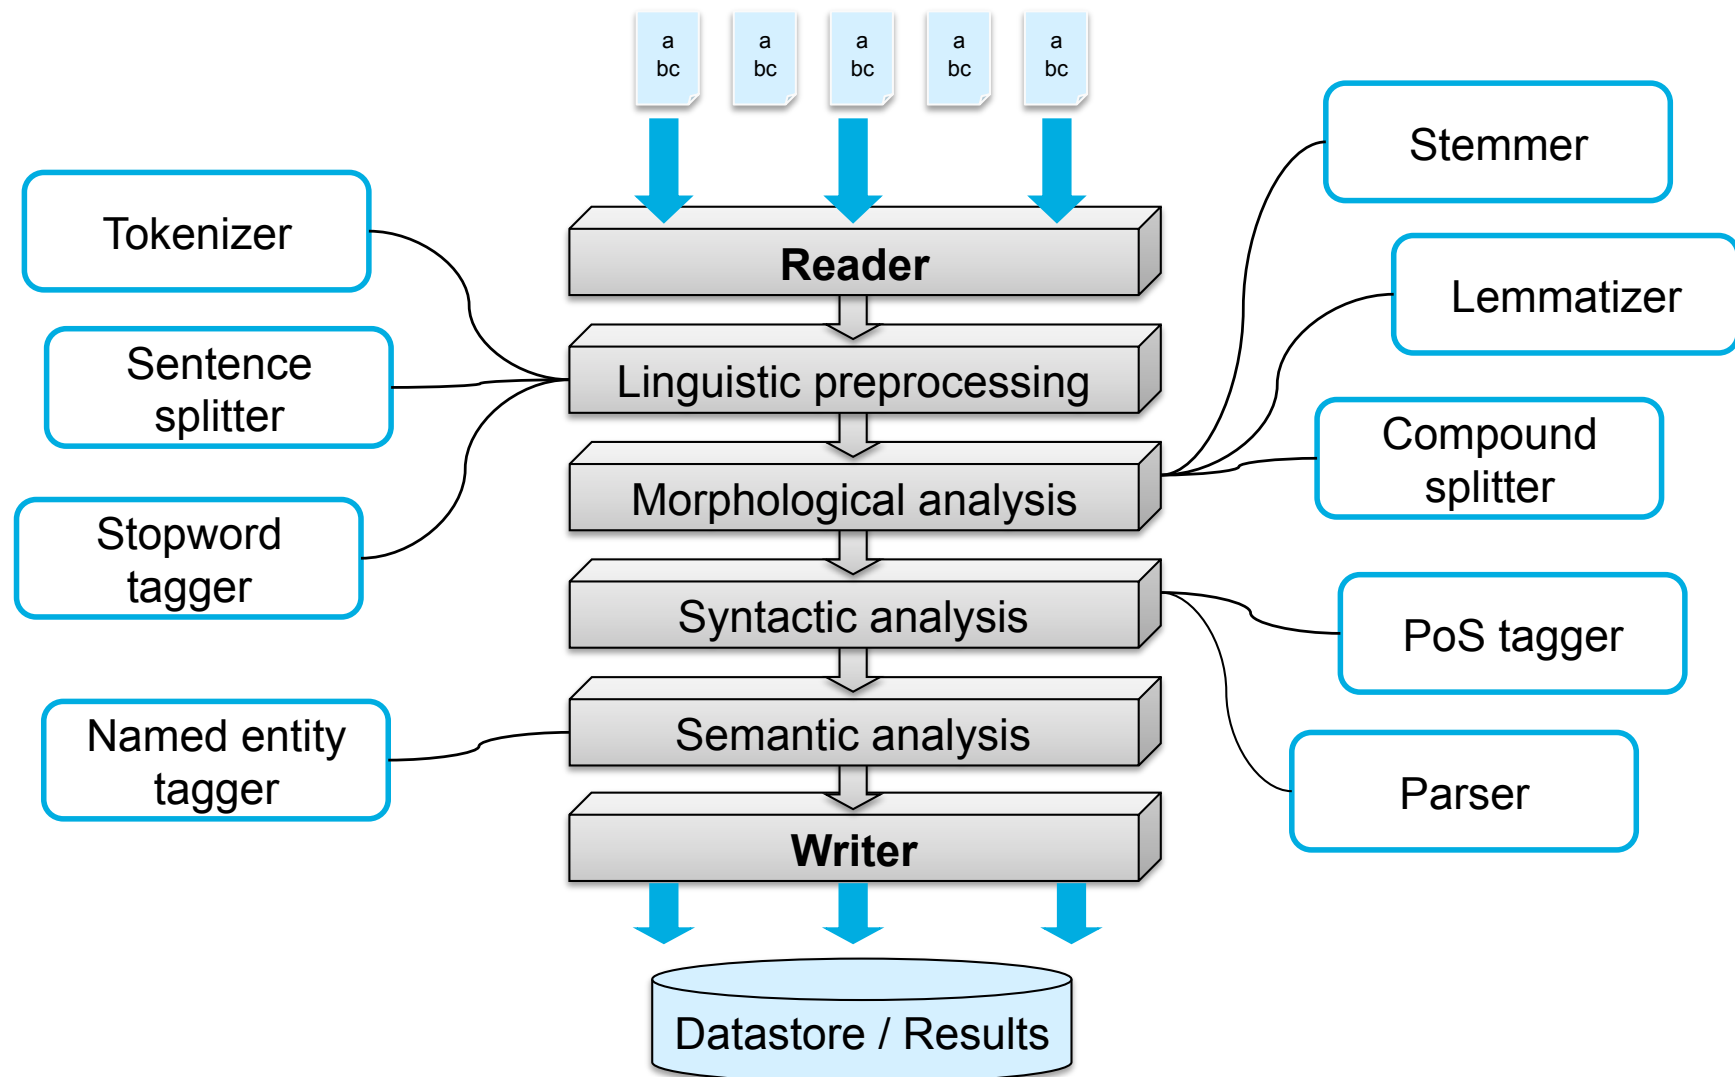


TECHNISCHE
UNIVERSITÄT
DARMSTADT

Pipeline Architecture



TECHNISCHE
UNIVERSITÄT
DARMSTADT



Component – Collection Reader



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Empty data structure (CAS) is passed to the reader
- Reader sets text (SofA) and meta-data (e.g. language)

Reader

CAS

SofA	Language:	Latin
	DocumentText:	Ubi est Cornelia? Subito Marcus vocat: „Ibi Cornelia est, ibi stat!“

Component – Analysis Engine



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- The structure is passed to one analysis engine (AE) after the other
- Each analysis engine derives a bit of structure and records it (Annotation)

Reader

Tokenizer

Name
Detector

CAS

SofA

Language: Latin
DocumentText: Ubi est Cornelia?
Subito Marcus vocat:
„Ibi Cornelia est, ibi stat!“

Token(0, 3) Token(4, 7) Token(8,16) ...
Name(8, 16) Name(25, 31) ...

Component – CAS Consumer



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Do something interesting with the analyzed data

Reader

Tokenizer

Name
Detector

Name
Lister

Word
Counter

CAS

SofA

Language:

Latin

DocumentText:

Ubi est Cornelia?

Subito Marcus vocat:

„Ibi Cornelia est, ibi stat!“

Token(0, 3) Token(4, 7) Token(8,16)...

Name(8, 16) Name(25, 31) ...

Cornelia
Marcus

11 words

8 unique words

Pipeline Example



TECHNISCHE
UNIVERSITÄT
DARMSTADT

```
SimplePipeline.runPipeline(  
    createReaderDescription(TextReader.class,  
        TextReader.PARAM_SOURCE_LOCATION, "texts/**/*.txt"  
        TextReader.PARAM_LANGUAGE, "en"),  
  
    createEngineDescription(OpenNlpSegmenter.class),  
    createEngineDescription(MatePosTagger.class),  
    createEngineDescription(ClearNlpLemmatizer.class),  
    createEngineDescription(BerkeleyParser.class,  
        BerkeleyParser.PARAM_WRITE_PENN_TREE, true),  
    createEngineDescription(StanfordNamedEntityRecognizer.class),  
  
    createEngineDescription(XmiWriter.class,  
        XmiWriter.PARAM_TARGET_LOCATION, "output",  
        XmiWriter.PARAM_TYPE_SYSTEM_FILE, "TypeSystem.xml"));
```

UIMA Data Structures



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Common Analysis System (CAS)



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Provides access to primary data
- Stores secondary data aka annotations
- Functions like an in-memory database
 - Annotation types are like “tables”
 - There are “indexes”



- UIMA specification is platform-independent
- Cannot rely on type system of implementation language (Java, C++)
- UIMA provides an “Object-oriented” type-system with
 - Type -> class
 - Feature -> class member
 - Feature Structure -> instance
 - Single inheritance
 - Sub-type polymorphism
 - no methods or encapsulation
- Primitive types: integer, float, boolean, string
- Built-in complex types: arrays, lists, Annotation
- Type-system forms communication contract between components

Java + CAS = JCas



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- JCas maps CAS types into the Java type system
- JCasGen generates Java classes from the XML type system descriptor
 - Token.java – feature structure wrapper with getters and setters
 - Token_type.java – type wrapper (cf. Java ‘Class’ class)
- JCas wrappers cannot be used stand-alone
- Type system descriptors still needed to initialize the underlying CAS

Code: CAS (UIMA)

File: CasAndJCasExample



TECHNISCHE
UNIVERSITÄT
DARMSTADT

```
TypeSystemDescription tsd = new TypeSystemDescription_impl();
TypeDescription tokenTypeDesc = tsd.addType("Token", "", CAS.TYPE_NAME_ANNOTATION);
tokenTypeDesc.addFeature("length", "", CAS.TYPE_NAME_INTEGER);

CAS cas = CasCreationUtils.createCas(tsd, null, null);
cas.setDocumentText("This is a test.");

Type tokenType = cas.getTypeSystem().getType("Token");
cas.addFsToIndexes(cas.createAnnotation(tokenType, 0, 4));
cas.addFsToIndexes(cas.createAnnotation(tokenType, 5, 7));
cas.addFsToIndexes(cas.createAnnotation(tokenType, 8, 9));
cas.addFsToIndexes(cas.createAnnotation(tokenType, 10, 14));
cas.addFsToIndexes(cas.createAnnotation(tokenType, 14, 15));

Feature lengthFeat = tokenType.getFeatureByBaseName("length");
AnnotationIndex<AnnotationFS> tokenIdx = cas.getAnnotationIndex(tokenType);
for (AnnotationFS token : tokenIdx) {
    token.setIntValue(lengthFeat, token.getCoveredText().length());
}

for (AnnotationFS token : tokenIdx) {
    System.out.println(token.getCoveredText() + " - " +
        token.getFeatureValueAsString(lengthFeat));
}
```

Code: JCas (uimaFIT)

File: CasAndJCasExample



TECHNISCHE
UNIVERSITÄT
DARMSTADT

```
JCas jcas = JCasFactory.createJCas();  
jcas.setDocumentText("This is a test.");
```

```
new Token(jcas, 0, 4).addToIndexes();  
new Token(jcas, 5, 7).addToIndexes();  
new Token(jcas, 8, 9).addToIndexes();  
new Token(jcas, 10, 14).addToIndexes();  
new Token(jcas, 14, 15).addToIndexes();
```

```
for (Token token : select(jcas, Token.class)) {  
    token.setLength(token.getCoveredText().length());  
}
```

```
for (Token token : select(jcas, Token.class)) {  
    System.out.println(token.getCoveredText()+" - "+token.getLength());  
}
```

Navigating the CAS with JCasUtil/CasUtil



- `select(cas, type)`
- `selectAll(cas)`
- `selectSingle(cas, type)`
- `selectSingleRelative(cas, type, n)`
- `selectBetween(type, annotation1, annotation2)`

```

// CAS version
Type tokenType = CasUtil.getType(cas, "my.Token");
for (AnnotationFS token : CasUtil.select(cas, tokenType)) {
    ...
}


// JCas version
for (Token token : JCasUtil.select(jcas, Token.class)) {
    ...
}
```

uimaFIT type system detection



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- uimaFIT example doesn't explicitly load/create type system
- Type system detection mechanism
- Types defined in XML descriptor files
- uimaFIT scans classpath for type system descriptor files



Type System Editor (Eclipse)

File: *typeSystemDescriptor.xml* (gsc12013-slides)



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Type System Definition

▼ **Types (or Classes)**

The following types (classes) are defined in this analysis engine descriptor.
The grayed out items are imported or merged from other descriptors, and cannot be edited here. (To edit them, edit their source files).

Type Name or Feature Name	SuperType or Range	Element Typ
<input type="checkbox"/> de.tudarmstadt.ukp.tutorial.gsc12013.slides.Token	uima.tcas.Annotation	
length	uima.cas.Integer	
de.tudarmstadt.ukp.tutorial.gsc12013.slides.Name	uima.tcas.Annotation	
de.tudarmstadt.ukp.tutorial.gsc12013.slides.Sentence	uima.tcas.Annotation	
de.tudarmstadt.ukp.tutorial.gsc12013.slides.Paragraph	uima.tcas.Annotation	

Add Type
Add...
Edit...
Remove
Export...
JCasGen
☐ limited

- JCasGen makes UIMA types available as Java Classes

In which order does select() return annotations?

- Type priorities define iteration order over features structures
- Example
 - A Sentence and a Paragraph start and end at the same position
- Which should be returned first?
 - Paragraph conceptually „larger“, but UIMA cannot know that by itself
- User specifies type priority: Paragraph, Sentence

Code: Type Priorities (JCas + uimaFIT)



TECHNISCHE
UNIVERSITÄT
DARMSTADT

File: TypePrioritiesExample

```
TypePriorities prio = createTypePriorities(  
    Paragraph.class, Sentence.class, Token.class);  
  
JCas jcas = createCas(createTypeSystemDescription(), prio,  
    null).getJCas();  
jcas.setDocumentText("This is a test.");  
  
new Token(jcas, 0, 4).addToIndexes();  
new Token(jcas, 5, 7).addToIndexes();  
new Token(jcas, 8, 9).addToIndexes();  
new Token(jcas, 10, 14).addToIndexes();  
new Token(jcas, 14, 15).addToIndexes();  
new Sentence(jcas, 0, 15).addToIndexes();  
new Paragraph(jcas, 0, 15).addToIndexes();  
  
for (Annotation a : select(jcas, Annotation.class)) {  
    System.out.println("[" + a.getType().getShortName() + "/" +  
        a.getBegin() + "-" + a.getEnd() + "]" + a.getCoveredText());  
}
```

Components

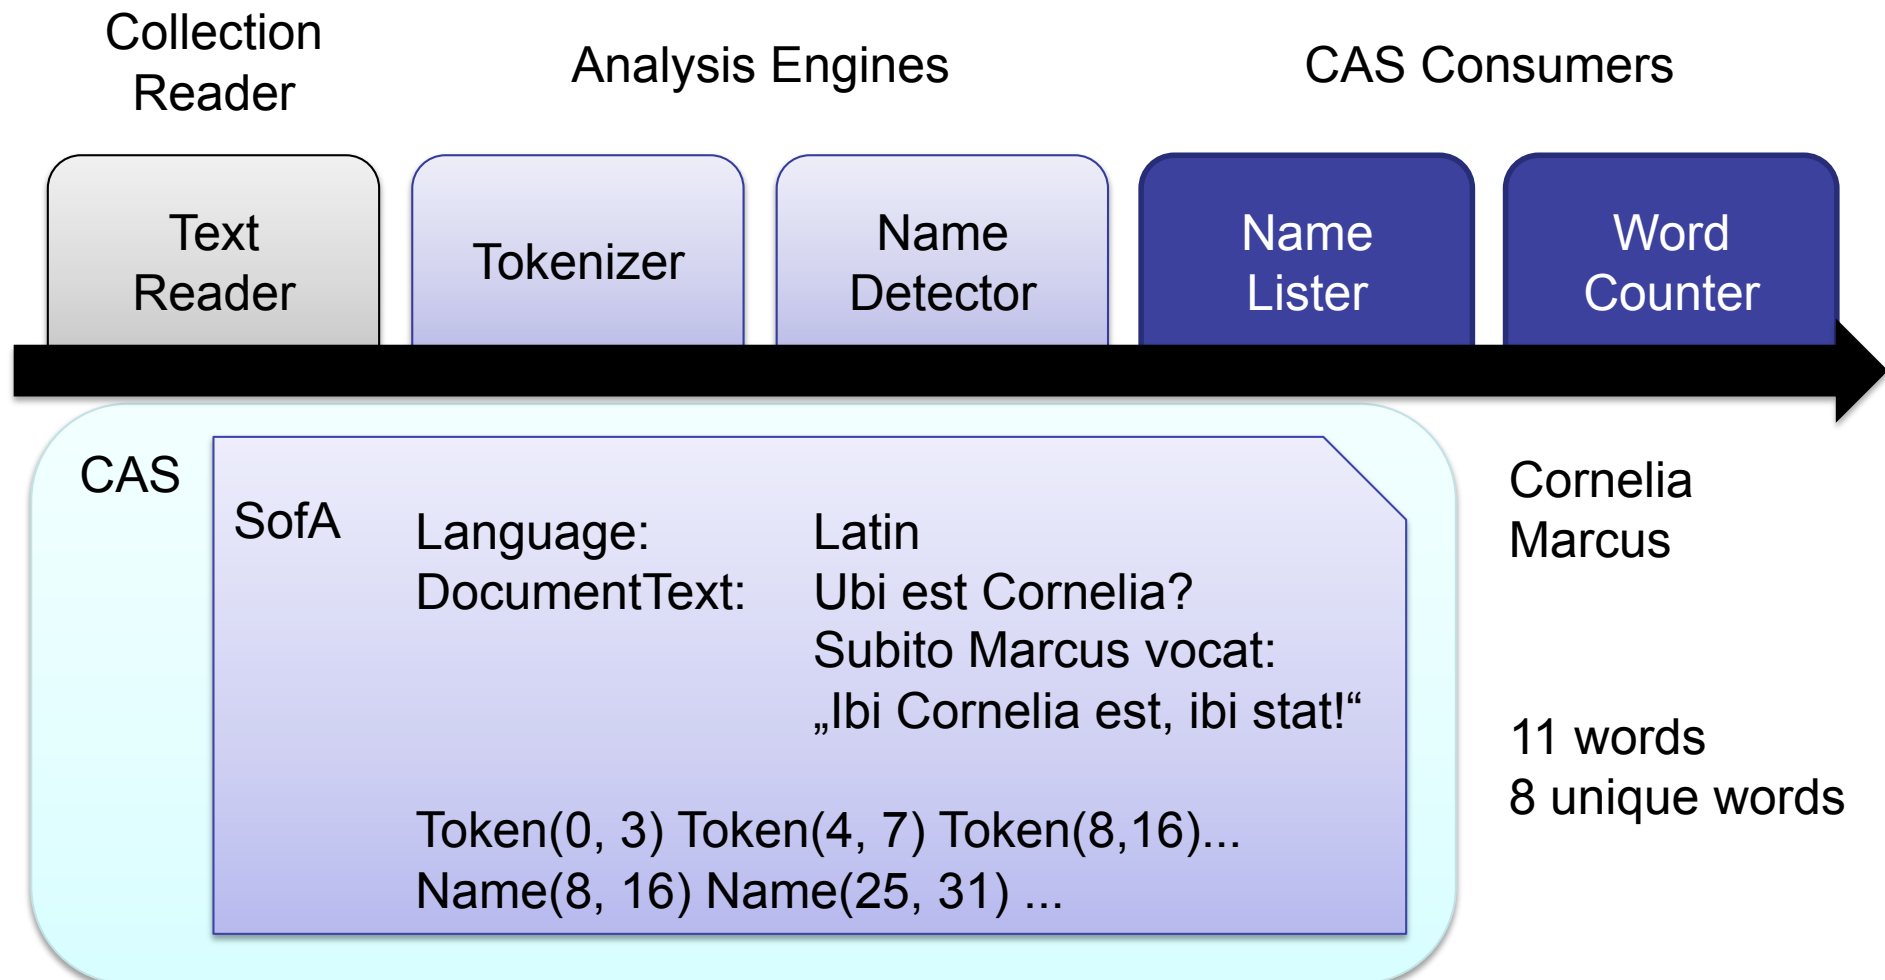


TECHNISCHE
UNIVERSITÄT
DARMSTADT

Components



TECHNISCHE
UNIVERSITÄT
DARMSTADT



API – Life-Cycle Events



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Component life-cycle events
 - initialize()
 - configure component from context
 - apply parameters, initialize resources
 - reconfigure()
 - context has changed
 - re-initialize component
 - destroy()
 - release resources
- Processing life-cycle events
 - collectionProcessComplete()
 - last CAS in collection has been processed
 - aggregated analysis complete
 - batchProcessComplete()
 - last CAS in batch has been processed
 - recoverable components set checkpoint
- Other
 - typeSystemInit()
 - type-system has changed

(CASAnnotator_ImplBase only)

API – Processing Methods



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- **CollectionReader**
 - hasNext() - another CAS can be filled
 - getNext() - fill the next CAS
 - getProgress() - report current progress to execution engine
- **AnalysisEngine**
 - process() - process/modify CAS
- **CasConsumer**
 - process() - process CAS; do not modify

Code: initialize() (uimaFIT)

File: TextFileReader



TECHNISCHE
UNIVERSITÄT
DARMSTADT

```
public static final String PARAM_PATH = "path";
@ConfigurationParameter(name = PARAM_PATH, mandatory = true)
protected File path;

public static final String PARAM_FILENAME_PATTERN = "filenamePattern";
@ConfigurationParameter(name = PARAM_FILENAME_PATTERN, mandatory = true,
    defaultValue = ".*\\.txt")
protected String filenamePattern;

public static final String PARAM_LANGUAGE = "language";
@ConfigurationParameter(name = PARAM_LANGUAGE, mandatory = true)
protected String language;

protected Queue<File> files;
protected int totalFiles;

public void initialize(UimaContext aContext) throws
ResourceInitializationException {
    super.initialize(aContext);
    files = new LinkedList<File>();
    collectFiles(path, filenamePattern, files);
    totalFiles = files.size();
}
```

Code: process() (uimaFIT)

File: NameAnnotator



TECHNISCHE
UNIVERSITÄT
DARMSTADT

```
public static final String PARAM_DICTIONARY_FILE = "dictionaryFile";
@ConfigurationParameter(name = PARAM_DICTIONARY_FILE, mandatory = true)
private File dictionaryFile;

private Set<String> names;

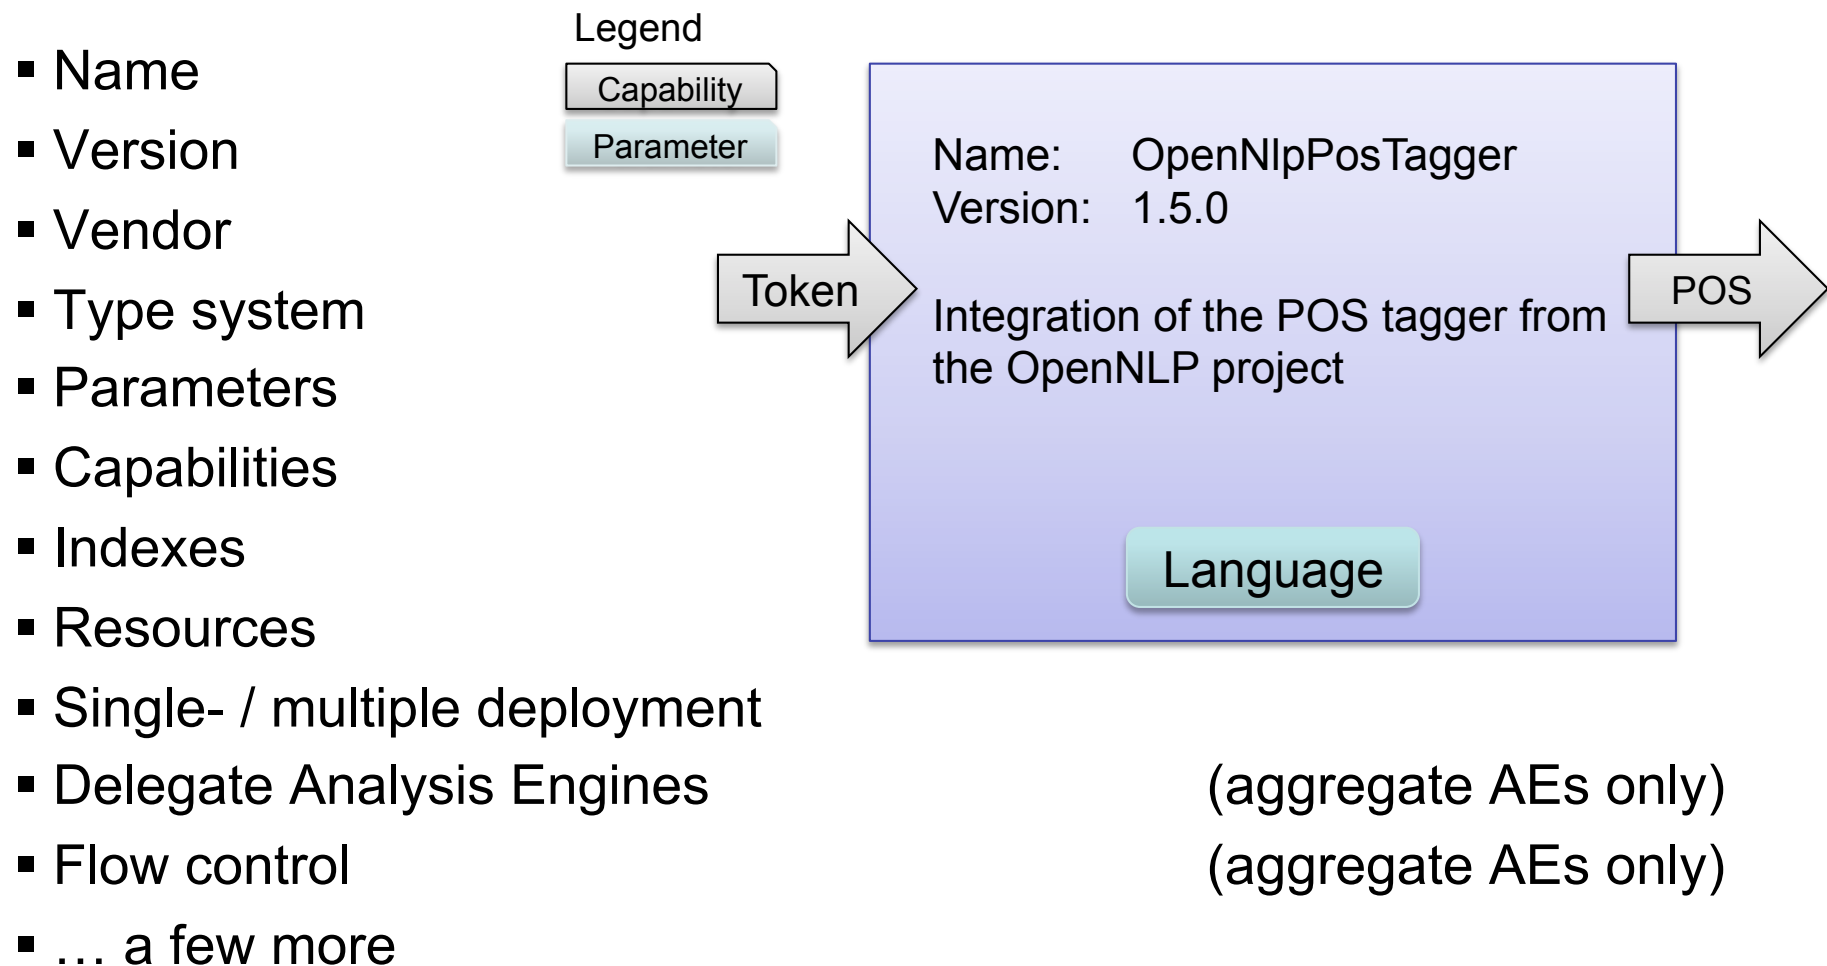
public void initialize(UimaContext aContext)
{
    super.initialize(aContext);
    names = new HashSet<String>(readLines(dictionaryFile));
}

public void process(JCas jcas)
{
    // Annotate tokens contained in the dictionary as name
    for (Token token : select(jcas, Token.class)) {
        if (names.contains(token.getCoveredText())) {
            new Name(jcas, token.getBegin(), token.getEnd()).addToIndexes();
        }
    }
}
```


Figure: Analysis Engine Descriptor



TECHNISCHE
UNIVERSITÄT
DARMSTADT



XML Descriptors – Pro & Contra



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Pro
 - “Officially preferred” form of configuring UIMA components/resources
 - Widely supported by UIMA tooling
 - XML elements usually correspond 1:1 to Java classes
- Contra
 - Mix declaration/documentation and configuration
 - Not included when refactoring code
 - No convenient API for use in Java (remedy: uimaFIT factories)

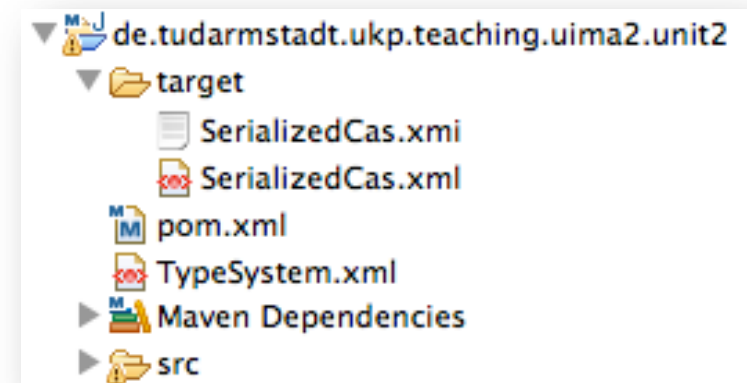
Persisting and loading a CAS



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Available serialization formats
 - XCAS – proprietary UIMA XML format (rarely used today)
 - XMI – standard XML representation of object graphs

- Type-system definition not included!



- Tip
 - Persist type system as “TypeSystem.xml” at project root
 - Open and XMI file in that project with the CAS Editor

Code: CAS to XMI (de)-serialization



TECHNISCHE
UNIVERSITÄT
DARMSTADT

File: CasPersistenceExample

```
CAS cas = createCas(createTypeSystemDescription(), null, null);  
populateCas(cas);  
  
FileOutputStream out = new FileOutputStream("target/SerializedCas.xmi");  
XmiCasSerializer.serialize(cas, out);  
closeQuietly(out);  
  
CAS loadedCas = createCas(createTypeSystemDescription(), null, null);  
FileInputStream in = new FileInputStream("target/SerializedCas.xmi");  
XmiCasDeserializer.deserialize(in, loadedCas);  
closeQuietly(in);  
  
createPrimitive(PrintConsumer.class).process(loadedCas);  
  
FileOutputStream typeOut = new FileOutputStream("TypeSystem.xml");  
createTypeSystemDescription().toXML(typeOut);  
closeQuietly(typeOut);
```

CAS Editor



TECHNISCHE
UNIVERSITÄT
DARMSTADT

The screenshot displays the CAS Editor interface. The main window is titled "SerializedCas.xml" and contains the text "This is a test." with the word "is" selected. To the right is the "FeatureStructure View" panel, which shows a list of tokens: "Token (id=13)", "Token (id=18)", "Token (id=23)", "Token (id=28)", and "Token (id=33)". The "Type" dropdown is set to "de.tudarmstadt.ukp.teach". Below the main editor are two "Edit View" panels, each showing a table of features and values.

Feature	Value
▶ sofa	[Sofa]
begin	5
end	7
length	0

Feature	Value
▶ sofa	[Sofa]
begin	5
end	7
length	0

example1.txt.xml

```

<frank> can you help me install GTA3?
<knightmare> first, shut down all programs you aren't using
frank has quit IRC. (Quit)
<knightmare> ...

```

Outline FeatureStructure View

Text

- CONJ
- ▼ ART
 - all
- ▼ ADJ
 - frank
 - first
 - frank
- ADV
- V
- N
- ▼ O
 - <
 - >
 - <
 - >
 - ,
 - (
 -)
 - <
 - >
 - CARD
 - ▼ NN
 - knightmare
 - programs
 - aren't
 - knightmare
 - ▼ NP
 - GTA3
 - IRC
 - Annotation

Edit View

Feature	Value
► sofa	[Sofa]
begin	86
end	92
PosValue	NN

Edit View

Feature	Value
► sofa	[Sofa]
begin	86
end	92
PosValue	NN

Annotation Styles

Type	Style
<input type="checkbox"/> Person	BACKG
<input type="checkbox"/> Sentence	BACKG
<input type="checkbox"/> Time	BACKG
<input type="checkbox"/> Token	BACKG
<input checked="" type="checkbox"/> ADJ	BACKG
<input checked="" type="checkbox"/> ADV	BACKG
<input checked="" type="checkbox"/> ART	BACKG
<input checked="" type="checkbox"/> CARD	BACKG
<input checked="" type="checkbox"/> CONJ	BACKG
<input checked="" type="checkbox"/> N	BACKG
<input checked="" type="checkbox"/> NN	BACKG
<input checked="" type="checkbox"/> NP	BACKG
<input checked="" type="checkbox"/> O	BACKG
<input type="checkbox"/> Animal	BACKG
<input type="checkbox"/> Cardinal	BACKG
<input type="checkbox"/> ContactInfo	BACKG
<input type="checkbox"/> Date	BACKG
<input type="checkbox"/> Disease	BACKG

DKPro Core Component Collection

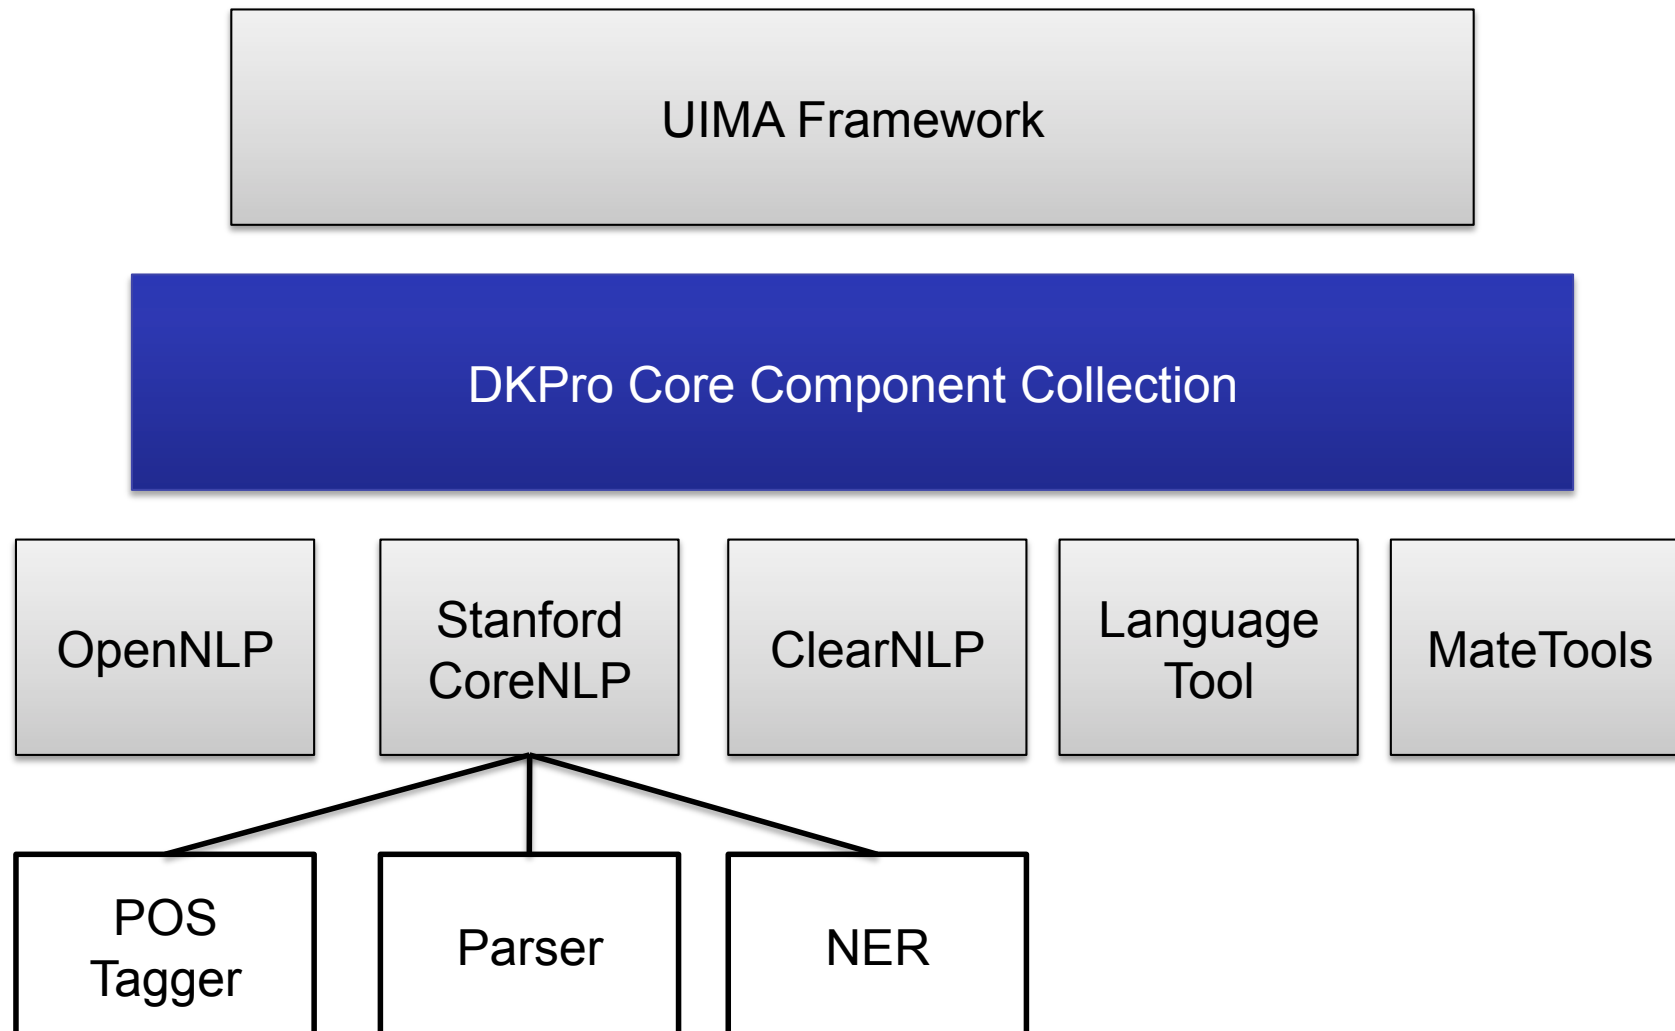


TECHNISCHE
UNIVERSITÄT
DARMSTADT

What's a component collection?



TECHNISCHE
UNIVERSITÄT
DARMSTADT



- DKPro Core is an integration framework
 - Processing: tools and models
 - Primary data: corpora
 - Auxiliary data: other language resources (e.g. lexical resources)
- Primarily integration of existing work, not original work
- Contribution of DKPro Core is the integration itself
- Open Source under Apache Software License & GNU Public License

<http://dkpro-core-asl.googlecode.com>

<http://dkpro-core-gpl.googlecode.com>

Stuff has to “just work”, everywhere.

- **Simplicity**

- Common data types used by all components
- Common set of parameters across components
- Sensible parameters defaults for minimal need for configuration
- Convenient deployment of components and resources
- Compose powerful pipelines with a few lines of code

- **Modularity**

- Use only what you need

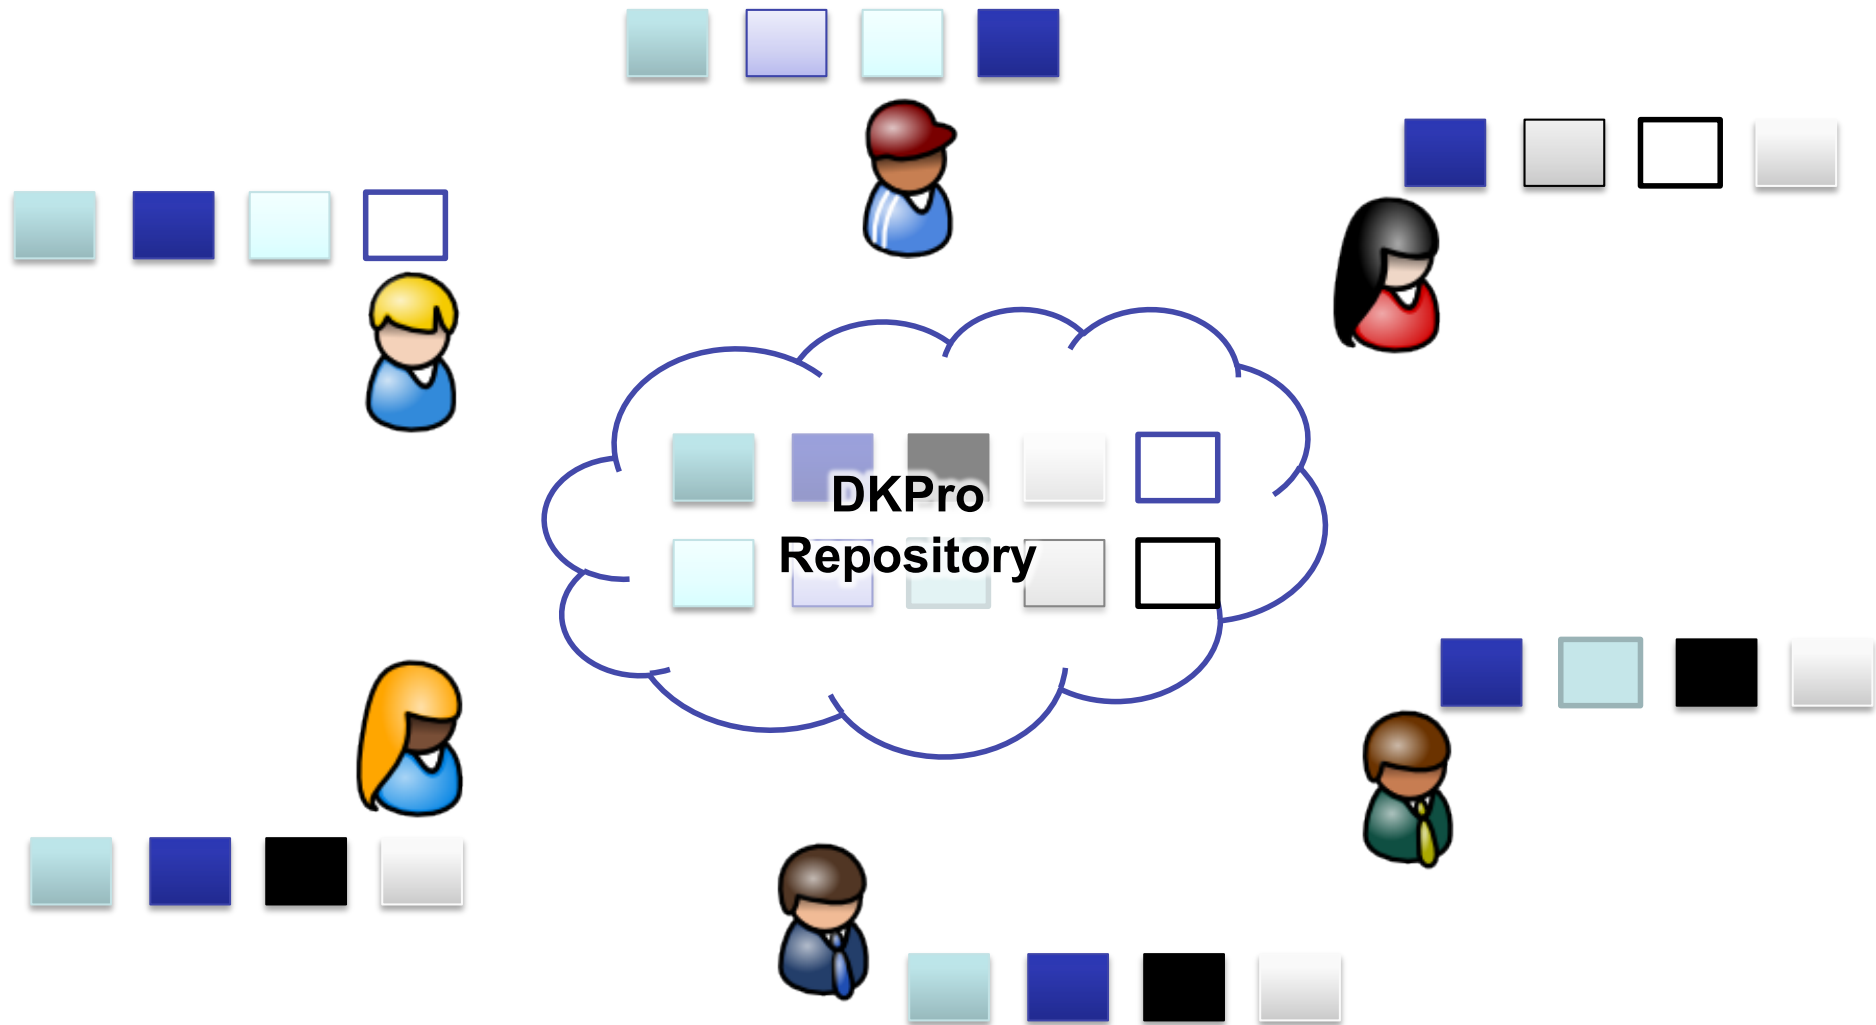
- **Flexibility**

- Override parameters for fine-grained control
- Extend data types with custom fields
- Customize type mappings

Managing Deployment



TECHNISCHE
UNIVERSITÄT
DARMSTADT



UKP OSS Component Repository

Publish component



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Overview

Artifact

Group Id:

Artifact Id*:

Version:

Packaging:

Parent

Group Id*:

Artifact Id*:

Version*:

Relative Path:

Project

Name:

URL:

Description:

Inception:

Organization

SCM

Issue Management

Continuous Integration

Properties



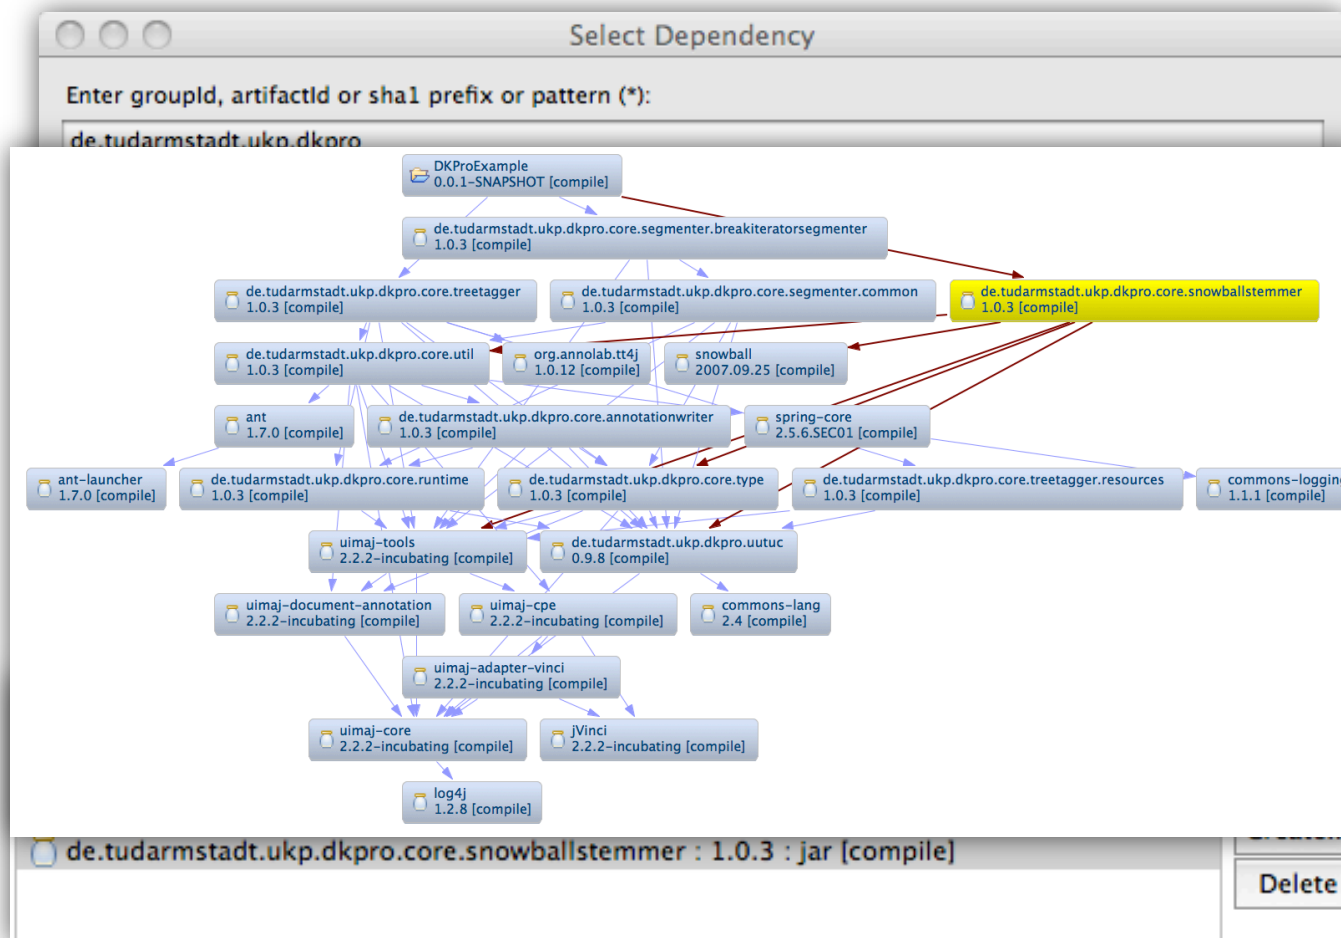

UKP OSS Component Repository

Retrieving components



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Component
Repository





Integrated Tools

- Stanford NLP
- OpenNLP
- Mate-Tools
- ClearNLP
- LanguageTool
- TreeTagger
- JWordSplitter
- Snowball Stemmer
- TextCat
- MaltParser
- MstParser
- BerkeleyParser
- MeCab
- Jazzy
- ...

Supported Formats

- Text
- PDF
- TIGER XML
- TEI XML
- BNC XML
- Negra Export
- SQL Databases
- Google web1t n-grams
- ...

Readers and Writers



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- **Common parameters**

- Source / target location
- Source / target encoding
- ANT-like patterns (for readers)
- Language (for readers)

- **Common features**

- Read data from file system, ZIP/JAR archives or classpath
- Preserve directory structure on write for recursive reads

Some currently supported corpora / resources



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- British National Corpus (BNC-XML)
- Wacky Corpora (OpenCWB format)
- TüBa D/Z (NEGRA export format)
- Tiger Corpus (Tiger XML)
- Digitale Bibliothek (TEI XML)
- Brown Corpus (TEI XML)
- ACL Anthology Reference Corpus (Text)
- ...

- Google Web1T n-grams
 - Can also easily build your own n-gram database with DKPro/jweb1t

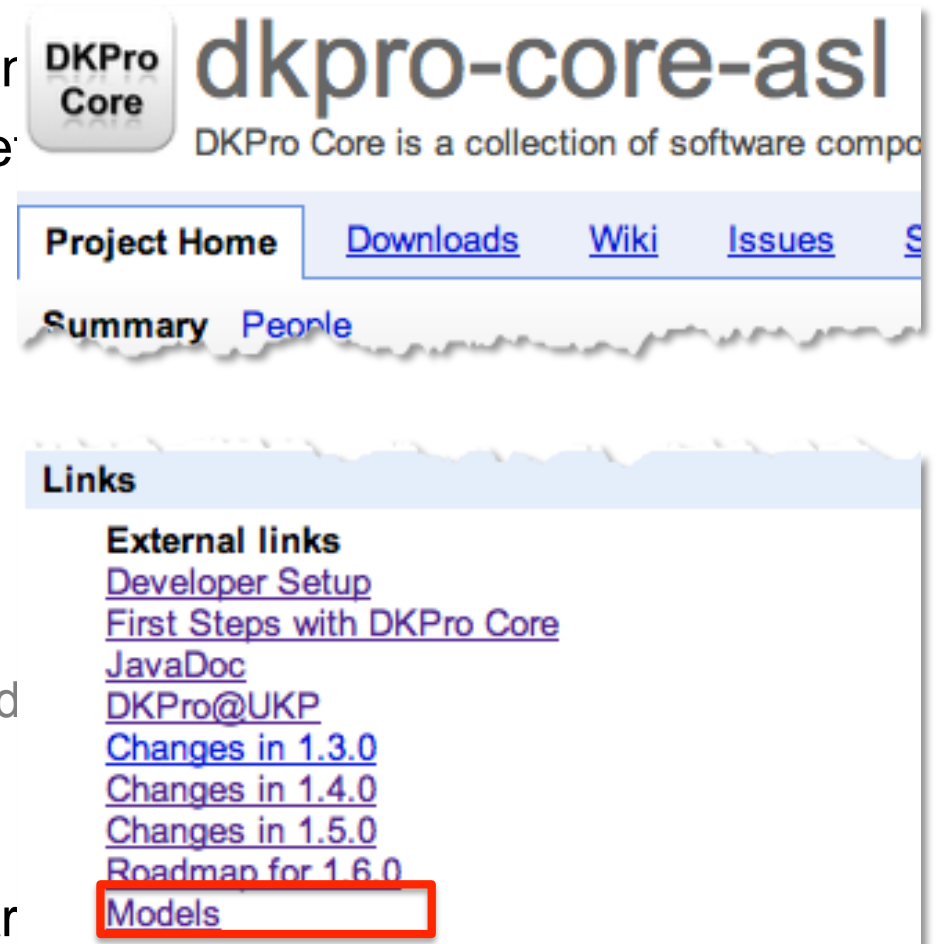
We are working on supporting more corpora and corpus formats

Good range of pre-trained models



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Upstream models packaged for conversion
- Package includes additional model measures
- 90+ models
- 20+ tools
- 15+ languages
- Best supported
 - English (Penn Treebank Tagset, Stanford)
 - German (STTS Tagset, Negra/Tiger)
- Unfortunately not all models are redistributed
- ... we are interested in creating/collecting more models



DKPro Type System Overview



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Meta Data

```
DocumentMetaData
<String documentTitle>
<String collectionId>
<String documentId>
<String documentUri>
<String documentBaseUri>
```

Segmentation

Document

Token

LinkingMorpheme

Heading

Compound

Ngram

Paragraph

Split

StopWord

Sentence

CompoundPart

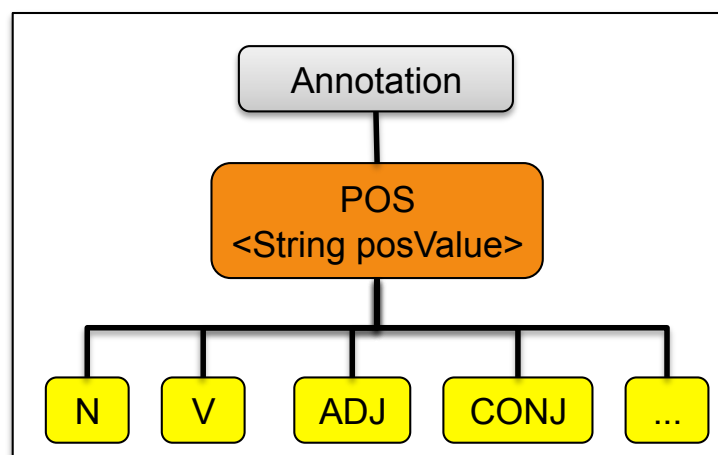
LexMorph

POS

Stem

Lemma

Morpheme



Syntax

Constituent

Dependency

Chunk

Coreference

CoreferenceChain

CoreferenceLink

Semantic Role Labeling

SemanticPredicate

SemanticArgument

Named Entities

NamedEntity

Person

Location

...etc...

UIMA type mappings

- Tags mapped to UIMA types (configurable)
- **Generic:** Original tags stored in a *value* feature, e.g. `POS.value`
- **Coarse Grained:** Currently supported for Part-of-Speech tags
 - 13 coarse grained part-of-speech tags
 - ADJ, ADV, ART, CARD, CONJ, N (NP, NN), O, PP, PR, V, PUNC
- Convenient coarse-grained processing across languages
- Similar “Universal Part-of-Speech” tag-set published @ LREC 2012
 - *Slav Petrov, Dipanjan Das and Ryan McDonald*
 - Defines mappings for 25 tagsets in 22 languages
 - Will be adopted in a future DKPro Core release
- Similar coarse-grained mappings may come for syntax, dependencies, ...

Analysis Engines



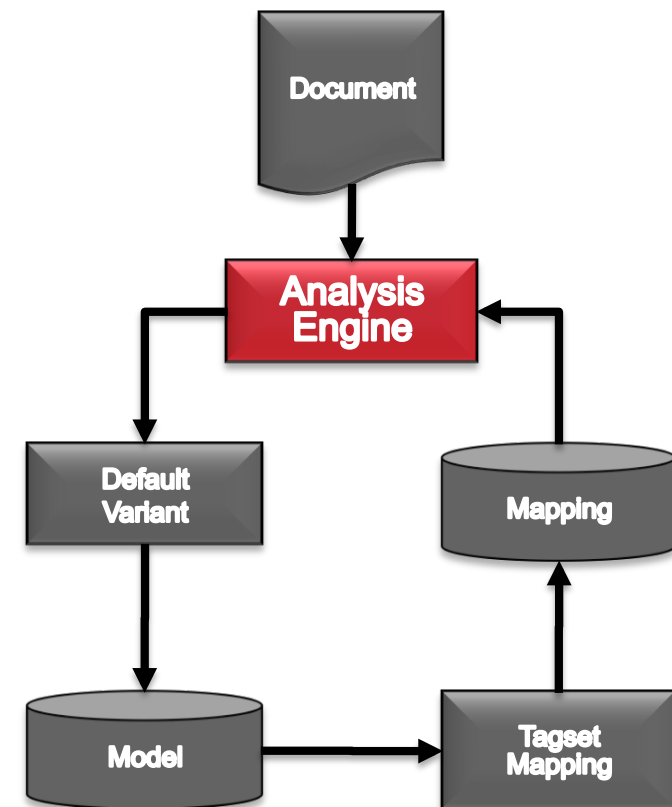
TECHNISCHE
UNIVERSITÄT
DARMSTADT

▪ Common parameters

- Model location
- Model encoding
- Model variant
- Mapping location
- Language

▪ Common features

- Load model depending on document language
- Print model tag set to log
- Default variants



classpath:/de/tudarmstadt/ukp/dkpro/core/opennlp/lib/tagger-**\${language}-\${variant}**.bin
classpath:/de/tudarmstadt/ukp/dkpro/core/api/lexmorph/tagset/**\${language}-\${pos.tagset}**-pos.map

DKPro Core and Groovy



TECHNISCHE
UNIVERSITÄT
DARMSTADT

```
#!/usr/bin/env groovy
@Grab(group='de.tudarmstadt.ukp.dkpro.core',
      module='de.tudarmstadt.ukp.dkpro.core.opennlp-asl',
      version='1.5.0')

import de.tudarmstadt.ukp.dkpro.core.opennlp.*;
import org.apache.uima.fit.factory.JCasFactory;
import org.apache.uima.fit.pipeline.SimplePipeline;
import de.tudarmstadt.ukp.dkpro.core.api.segmentation.type.*;
import de.tudarmstadt.ukp.dkpro.core.api.syntax.type.*;
import static org.apache.uima.fit.util.JCasUtil.*;
import static org.apache.uima.fit.factory.AnalysisEngineFactory.*;

def jcas = JCasFactory.createJCas();
jcas.documentText = "This is a test";
jcas.documentLanguage = "en";

SimplePipeline.runPipeline(jcas,
    createEngineDescription(OpenNlpSegmenter),
    createEngineDescription(OpenNlpPosTagger),
    createEngineDescription(OpenNlpParser,
        OpenNlpParser.PARAM_WRITE_PENN_TREE, true));

select(jcas, Token).each { println "${it.coveredText} ${it.pos.posValue}" }
select(jcas, PennTree).each { println it.pennTree }
```

Hands-on



TECHNISCHE
UNIVERSITÄT
DARMSTADT

<https://dl.dropboxusercontent.com/u/11205710/iuahsdfhksad/ws-tutorial.zip>

Example projects



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- **gscl2013-types**
 - Defining and using your own annotation types
- **gscl2013-pipeline**
 - Four ways of building a UIMA pipeline
 1. No reader, no writer
 2. No reader, with writer
 3. With reader, no writer
 4. With reader, with writer
 - Simple interactive analysis pipeline
 - Extended interactive analysis pipeline with language detection
- **gscl2013-dkpro**
 - DKPro Core monster pipeline wildly mixing components from different vendors
- **gscl2013-pos-ensemble**
 - Building an improved part-of-speech tagger using an ensemble of taggers
- **gscl-2013-ruta**
 - Running Ruta within an uimaFIT pipeline